




Agentic Architecture in Practice

Harness · LLM · Directives

A working model for how agentic systems are designed, governed, and extended — with **RockBot** as the example.

Rockford Lhotka · rockbot.dev · blog.lhotka.net



THE THESIS

What **is** an agent?

Not "an LLM with tools bolted on." In practice it's three distinct parts:



Harness

The code that controls execution
— the loop.

LLM

Reasons and generates. A black
box.

Directives


Shape behavior — scope, tone,
rules.

Pull any one out and you don't have an agent.



The LLM is a **black box**

It's a microservice: **text in** → **text out**. Nothing more.



It doesn't remember the last call.

It can't call a tool — it just **writes** that it wants to.


It takes no action. It only generates text.

So where does all the "agent" behavior come from? →

PART 2 OF 3 — THE HARNESS

The harness is **the loop**

This is where the real engineering lives.



The **inner loop** — call a tool, feed the result back, call the LLM again — is what makes an agent feel agentic.

Directives **shape behavior**

Scope · tone · rules-for-success. In RockBot, literally markdown files, composed into the system prompt.

soul.md

Scope & identity.

Who the agent is, what it's for.

directives.md

Rules.

Constraints that keep it safe and on-task.

style.md

Tone.

How it communicates.

Layered & composable — base persona, then per-task, then per-skill:

soul + directives + style

+

task directive

+

skill

→

system prompt

Change behavior without changing code.

THE CONCRETE EXAMPLE

Meet RockBot

An **event-driven, process-isolated** autonomous agent framework for .NET.

"Nothing trusts
the LLM."

message bus (RabbitMQ)

isolated processes

least privilege

MIT · open source


Everything that follows is a real subsystem in this codebase.



REACHING OUTWARD

Tools: how the harness **acts**

The LLM only **asks** for a tool. The harness routes and runs it — then feeds the result back.




RockBot leans on **MCP** heavily — secrets live in the MCP servers, not the agent. (more on that later)



REACHING OUTWARD

Peers: agents calling agents

The **A2A** protocol lets RockBot delegate to other agents — or be delegated to.



A tool call reaches a service. An **A2A** call reaches a reasoning peer — its own harness, LLM, and directives.

THE PIVOT

The harness handles a lot

Once the core (loop + LLM + directives) is solid, this is what becomes possible:



The rest of the talk walks the spokes.

Memory: three tiers + a graph



Conversation

Last N turns. Session-scoped, ephemeral. The here-and-now.



Working

Shared, namespaced scratch space with a TTL. Hand-off between sessions, tasks & subagents.



Long-term

Permanent facts & preferences. Recalled per-turn by BM25 (+ vector) search.



Knowledge graph

Entities & relationships for structured, relational reasoning.

The trick isn't storing it — it's **injecting only what's relevant**, every turn, without token bloat.

Delta injection: a memory recalled once isn't re-injected the same session.

Skills: **fixed** guides & **evolving** know-how

Memory stores **facts**. Skills store **how to do something** — markdown procedures, recalled by BM25.

Subsystem guides

Authored, fixed. "How to use the email MCP," "how working memory works." Ship with the agent.

Mutable skills


The agent **writes its own** — captures a procedure that worked, then refines it over time.



The refine step happens offline — in the dream. →

Tasks on a schedule

No user required. Cron-driven "patrol" tasks fire on their own — **each with its own directive.**



Same harness, same LLM — a **different directive** turns it into a different worker.

RockBot dreams

A background cycle that runs offline and **refactors the agent's own knowledge** — no user, no goals changed.

Consolidate memory

Merge duplicates, decay stale facts, mine anti-patterns from corrections.

Optimize skills

Refine, merge, and cluster the skills it has written.

Tune LLM routing

Send the right work to the right model — cost & quality.

Infer preferences

Notice patterns across conversations; learn how you like to work.

Yesterday's experience makes tomorrow's agent sharper — automatically.

Principle of least privilege

The core philosophy, restated: **nothing trusts the LLM**. The core agent holds **no keys, no passwords**.



A compromised prompt can't leak a secret the agent never had.

Untrusted code runs **sandboxed**

RockBot can run Python and pull web pages — but never in-process.

Python execution

Runs in an **ephemeral, low-privilege container**. Spun up per run, destroyed after. No host access, no secrets, minimal blast radius.

Web search & fetch

Search the web and retrieve pages — but that content is **untrusted input**, handled with the same suspicion as the LLM's output.

Same pattern everywhere: **assume the LLM (and the web) can be wrong or hostile** — and contain it.

THE WORKING MODEL



Get the core right, and it extends:



From abstract "agent" talk to a working model for how agentic systems are **designed, governed, and extended.**



rockbot.dev · github.com/MarimerLLC/rockbot · blog.lhotka.net